

Responsible-AI Enablement — Training Module Script

"Using Generative AI Responsibly at the Firm"

Format: Self-paced module, built in Articulate Rise **Audience:** Attorneys, associates, paralegals, and business professionals adopting an approved GenAI assistant **Length:** ~20 minutes · three chapters **Goal:** Confident, safe, repeatable use — prompt well, verify always, protect confidentiality — with attorney oversight throughout.

Independent work sample created for this application. Not affiliated with or endorsed by Winston Taylor. Use only firm-approved tools; never enter client or privileged data into an unapproved tool.

How to use this script with Articulate Rise's AI assistant

Paste one chapter at a time into Rise's AI builder. Each gives the learning objective, on-screen blocks (type in [brackets]), copy, and answer-keyed checks. Set the Rise theme to a light/cream background, near-black headings, and a single purple accent to match the firm brand.

INTRO — Why this module exists

[Cover / statement block] Headline: *AI helps the firm only if we use it well — and safely.* Subhead: Three habits make that true: prompt well, verify always, protect confidentiality. An attorney always owns the result.

[Text block] Generative AI can speed up drafting, research, summarization, and analysis. It can also hallucinate, and it can put confidential data at risk if used carelessly. This module builds the habits that let you get the upside without the downside — the same habits the firm's guardrails are built around.

[Knowledge check — multiple choice] Q: What is AI's role in your work product? - A) The final word, to save time - B) A first draft that an attorney reviews and owns ☐ - C) A replacement for legal judgment - D) A place to store client data Feedback (correct): AI supports judgment; a human reviews and owns every output.

CHAPTER 1 — Prompt well

[Statement block] *Objective:* Write prompts that produce a useful, defensible first draft.

[Numbered list block] 1. Give the model a clear role and task ("You are helping draft a research memo on X..."). 2. Provide the context it needs — but keep it non-confidential unless you're in an approved, firm-contained tool. 3. Specify format, audience, and constraints. 4. Iterate: refine the prompt instead of accepting the first answer.

[Knowledge check — multiple choice] Q: You need a summary of a sensitive client document. What's the safe move? - A) Paste it into any chatbot to save time - B) Use only an approved, firm-contained tool — or de-identify first ☐ - C) Email it to yourself first
Feedback (correct): Sensitive data only goes into approved tools. When unsure, treat it as confidential.

CHAPTER 2 — Verify always

[Statement block] *Objective:* Confirm every fact and citation before you rely on it.

[Numbered list block] 1. Treat every fact and citation as unverified until you check it. 2. Open the source — does the cited authority exist and say what's claimed? 3. Watch for confident-sounding but wrong details (hallucinations). 4. If it can't be verified, it doesn't go in the work product.

[Flashcard grid block] - **Fake or wrong citation** → open it; if it doesn't exist or doesn't support the point, cut it. - **Plausible but wrong fact** → cross-check against a trusted source. - **Over-confident tone** → confidence is not accuracy; verify anyway.

[Knowledge check — multiple response] Q: Which are part of verifying? (Select all) - Open and read the cited source ☐ - Cross-check key facts ☐ - Assume it's right because it sounds confident ☐ - Remove anything you can't verify ☐ Feedback: Verify the high-stakes points; never rely on unverified output.

CHAPTER 3 — Protect confidentiality (and keep humans in the loop)

[Statement block] *Objective:* Keep client and privileged data safe, and know when to

escalate.

[Numbered list block] 1. Know your data classification: what may be used with which tool. 2. Client and privileged data stays in approved, firm-contained tools only. 3. An attorney reviews and owns every AI-assisted output before it's used. 4. Know the escalation route for questions, suspected errors, or misuse — and use it early.

[Statement block]

When in doubt, treat it as confidential and ask. A quick check beats an incident.

[Knowledge check — scenario] Q: The tool returns a great paragraph but cites a case you can't find, and you used a public tool with a redacted doc. What do you do? - A) Use it — it reads well - B) Remove the unverifiable citation, and confirm the tool and data handling were approved ☐ - C) Send it to the client Feedback (correct): Cut what you can't verify, and check that both the tool and the data were within the guardrails.

CLOSE — wrap, job aid, final check

[Text block] You now have the three habits that make AI both useful and safe at the firm: prompt well, verify always, protect confidentiality — with an attorney owning the result.

[Download / file block] Attach the one-page **Responsible-AI quick card** (below).

[Final knowledge check — multiple choice] Q: What's the through-line of this module? - A) Trust AI fully to save time - B) Prompt well, verify always, protect confidentiality — with human oversight ☐ - C) Avoid AI entirely Feedback (correct): Useful and safe, together.

[Closing statement block] Headline: *Fast, and safe.* Subhead: More time for judgment and advocacy — with the firm and its clients protected.

One-page quick card (attach in the Close)

Responsible-AI quick card 1. **Prompt:** clear role/task/context; keep inputs non-confidential unless in an approved tool. 2. **Verify:** open every citation; cross-check key facts; cut what you can't confirm. 3. **Protect:** approved tools only for client/privileged data; know your data classification. 4. **Own it:** an attorney reviews every output; escalate questions, errors, or misuse early. *When in doubt, treat it as confidential and ask.*